
Lab: Deploy Azure Virtual Machine Scale Set (VMSS)

Lab Goal

Deploy a VMSS with:

- Ubuntu Linux VMs
 - A custom startup script to install a web server (Apache)
 - Azure Load Balancer
 - Autoscale rules (based on CPU)
-

Part 1: Create a Virtual Machine Scale Set

1. Log in to the Azure Portal

Go to: <https://portal.azure.com>

2. Search and open "Virtual Machine Scale Sets"

- Click "Create" > "Virtual machine scale set"
-

3. Basics Tab

Field	Value
Subscription	Your active subscription
Resource Group	VMSSLabRG (or create a new one)
Name	MyVMSS
Region	Choose a region with Availability Zones (e.g., East US)

Orchestration Mode	Uniform (default)
Image	Ubuntu Server 20.04 LTS
Authentication Type	Password or SSH public key
Username	azureuser
Password / SSH Key	Provide credentials

4. Instances Tab

Field	Value
Initial instance count	2
Availability Zone	Zone redundant (recommended)
VM Size	Standard B1s (cost-effective)

5. Disks Tab

- Use the default OS disk settings.
-

6. Networking Tab

Field	Value
Virtual network	Create new or use existing
Public IP address	Enable
Load balancing	Select "Azure Load Balancer"
Inbound ports	Allow HTTP (80) and SSH (22)

7. Management Tab

- Enable **Boot diagnostics** (optional)
 - Leave the rest as default
-

8. Health Monitoring (Optional)

- Can add a **health probe** for HTTP (port 80)
-

9. Scaling Tab

- Enable **Autoscaling**
 - Configure like:
 - **Minimum:** 2 instances
 - **Maximum:** 5 instances
 - **CPU threshold:** 70% (scale out) and 30% (scale in)
-

10. Advanced Tab

- Go to **"Custom data"** (cloud-init script)

Paste the following to install a web server and a test page:

```
#!/bin/bash
sudo apt-get update
sudo apt-get install -y apache2
echo "<h1>Welcome to Azure VM Scale Set!</h1>" | sudo tee /var/www/html/index.html
sudo systemctl start apache2
sudo systemctl enable apache2
```

11. Review + Create

- Review all settings and click **"Create"**
 - Deployment takes 3–5 minutes
-

Part 2: Test the Web App

1. Go to your VMSS Overview

- Copy the **Public IP** from the **Load Balancer**

2. Open a browser

Visit: `http://<public-ip>`

You should see:

 **Welcome to Azure VM Scale Set!**

Part 3: Trigger Autoscaling (Optional)

You can simulate CPU load to test scaling (on a VM instance):

```
sudo apt install stress -y
stress --cpu 1 --timeout 300
```

This should cause the scale set to spin up additional instances once CPU exceeds 70%.





Clean Up (To Avoid Charges)

Delete the **resource group**: VMSSLabRG

```
az group delete --name VMSSLabRG --yes --no-wait
```

-

Lab Summary

Task	Done
Deploy VMSS	
Install web app via startup script	
Configure Load Balancer	
Setup autoscaling	
Access public site	